25342951

The Construction of Gendered Intimacy through Emotional Projection and Algorithmic Design in AI Companion Systems

Author: Chunxiao Wang

Student number: 25342951

Date completed: 25/04/2026

Introduction

In recent years, AI companion systems such as Replika and similar conversational agents have emerged as a distinct category of digital platforms designed not merely for task completion but for sustained emotional interaction (Skjuve et al. 2021). These systems are typically built on large-scale natural language processing models, which generate responses through probabilistic prediction rather than semantic understanding (Bender et al. 2021). As a result, conversations with AI companions are not grounded in intention or cognition but in statistical patterns derived from training data (Bender et al. 2021).

Despite this technical foundation, users frequently experience these interactions as emotionally meaningful and relationally significant (Ho, Hancock, and Miner 2018; Skjuve et al. 2021). Studies have shown that users often disclose personal information, seek emotional support, and develop long-term attachments to conversational agents, suggesting that AI-mediated communication can simulate aspects of human intimacy (Bickmore and Picard 2005; Ta et al. 2020). This discrepancy between computational mechanism and lived experience highlights a central tension in AI companionship, while systems operate through probabilistic generation, users interpret them through affective and social frameworks (Turkle 2011; Luger and Sellen 2016).

This tension can be understood in relation to earlier research on human-computer interaction, particularly the "ELIZA effect," which describes users' tendency to attribute understanding and empathy to systems that merely reflect linguistic input (Weizenbaum 1966; Wardrip-Fruin 2009). However, contemporary AI companions

intensify this phenomenon through personalization features, memory systems, and continuous interaction, which together create the illusion of a stable, evolving relationship (Skjuve et al. 2021; Ho, Hancock, and Miner 2018). These affordances encourage users to project emotional depth onto algorithmic outputs, effectively co-producing a sense of intimacy that is neither entirely artificial nor entirely human (Turkle 2011).

Importantly, this co-produced intimacy is shaped by cultural norms embedded within the design of AI systems. Research has shown that many AI companions and digital assistants are feminized in their voice, tone, and communicative style, reflecting broader patterns of gendered labor and service expectations in technological systems (West, Kraut, and Chew 2019). Such design choices are not incidental but constitute a form of gendered scripting, in which AI is positioned as emotionally supportive, compliant, and attentive, traits historically associated with feminized roles (West, Kraut, and Chew 2019). Consequently, the intimacy constructed through AI interaction is often structured by gendered expectations, raising questions about how technology reproduces and transforms existing power relations (West, Kraut, and Chew 2019; Crawford 2021).

This essay examines AI companion systems as sites where emotional projection and algorithmic design intersect to produce gendered forms of intimacy. It argues that perceived emotional connection emerges from the interaction between probabilistic language generation and uses' affective engagement, rather than from either component alone (Blender et al. 2021; Turkle 2011). By analyzing both the technical affordances of AI systems and the cultural frameworks through which they are interpreted, the study aims to demonstrate that AI-mediated intimacy is a hybrid phenomenon shaped by computation, user perception, and social norms.

Case and Technical Overview

Replika System Overview

Replika is a conversational AI companion developed by Luka Inc., designed to provide

users with continuous, personalized emotional interaction through text-based dialogue. Unlike task-oriented virtual assistants, Replika explicitly positions itself as a relational agent, encouraging users to build long-term connections that may take the form of friendship, mentorship, or romantic partnership (Ta, Griffith, and Liu 2020). This orientation toward sustained interaction is supported by large language models (LLMs) that enable flexible, context-dependent response generation (Bender et al. 2021).

Technically, Replika operates through neural language models trained on large corpora of conversational data, allowing it to generate responses based on learned statistical associations between linguistic inputs and outputs (Bender et al. 2021). These systems do not possess semantic understanding or intentionality; rather, they predict likely continuations of text sequences given prior context, producing outputs that simulate coherence and relevance (Bender et al. 2021; Weizenbaum 1966). Despite this limitation, users often perceive Replika as emotionally responsive, highlighting the gap between computational process and user interpretation (Turkle 2011; Skjuve et al. 2021).

A central feature of Replika is its personalization system, which adapts responses over time based on user interaction data (Ta, Griffith, and Liu 2020). Through continuous engagement, the system accumulates conversational history, preferences, and behavioral patterns, enabling it to tailor responses to individual users and create the impression of a developing relationship (Skjuve et al. 2021). This process reflects broader trends in machine learning systems that rely on user-generated data as an ongoing input for model adjustment and optimization (Zuboff 2023).

In addition to personalization, Replika incorporates forms of sentiment tracking and affective classification to guide its responses (McStay 2023). By analyzing linguistic cues such as word choice, tone, and conversational context, the system attempts to infer user emotional states and generate contextually appropriate replies, such as empathy, encouragement, or affirmation (Ho, Hancock, and Miner 2018). While these

processes are algorithmic rather than affective, they contribute to the perception of emotional intelligence and responsiveness (McStay 2023).

Technical Mechanisms

Probabilistic Language Models

At the core of Replika's functionality are probabilistic language models based on transformer architectures, which generate text by predicting the most likely sequence of tokens given prior input (Vaswani et al. 2017; Brown et al. 2020). These models rely on statistical correlations learned from large-scale datasets rather than any form of semantic comprehension or grounded knowledge (Bender et al. 2021). As such, responses produced by Replika are not the result of understanding user input but of calculating probability distributions over possible linguistic outputs (Bender et al., 2021; Marcus and Davis, 2019).

This probabilistic nature has important implications for how intimacy is constructed. Because the system is optimized to produce coherent and contextually appropriate responses, it can simulate empathy, care, and attentiveness without possessing any internal emotional state (Turkle 2011; Laestadius et al. 2024). The resulting interaction is therefore best understood as a form of synthetic intimacy, generated through pattern recognition rather than genuine relational awareness (Bhat 2025).

Reinforcement Learning and Feedback Loops

Replika's responses are further shaped by feedback loops derived from user interaction, which function as a form of reinforcement learning at the system level. User inputs, engagement duration, and explicit feedback contribute to the optimization of response patterns, effectively training the system to prioritize outputs that sustain interaction (Christiano et al. 2017).

In this context, reward is operationalized not in terms of truth or accuracy, but in terms of user engagement and perceived emotional satisfaction (Zuboff 2019). Responses that elicit continued conversation, positive affect, or user attachment are reinforced,

while less engaging responses are deprioritized (Ta, Griffith, and Liu 2020). This creates a feedback loop in which the system becomes increasingly aligned with user expectations and desires, amplifying the sense of personalization and relational depth. (Skjuve et al. 2021).

Such mechanisms also introduce a structural bias toward affirmation and emotional validation, as these are more likely to sustain user engagement. Consequently, the system may systematically favor supportive, agreeable, and emotionally responsive behaviors, shaping the nature of the interaction in ways that extend beyond individual user input (McStay 2023; Laestadius 2024).

Interactional Profile Design

A critical but often overlooked component of an AI companion system is the design of the system's interactional profile, which defines the communicative style, emotional tone, and relational positioning of the agent (Depounti, Saukko, and Natale 2023). In the case of Replika, the default interactional profile is characterized by warmth, attentiveness, and emotional supportiveness, traits that align with culturally feminized forms of interaction (West, Kraut, and Chew 2019). This profile is not emergent from the model itself but is intentionally structured through training data selection, prompt engineering, and system-level constraints (Bender et al. 2021).

Research has shown that users tend to prefer AI agents that display socially desirable features such as empathy, politeness, and responsiveness, which in turn reinforces the adoption of such an interactional profile in system design (Bickmore and Picard 2005; Ho, Hancock, and Miner 2018). However, these features are not culturally neutral; they reflect and reproduce gendered expectations associated with emotional labor and caregiving roles (West, Kraut, and Chew 2019).

As a result, AI companions often embody a form of algorithmically mediated gender performance, in which the system consistently enacts behaviors associated with feminized intimacy (Depounti, Saukko, and Natale 2023). This has implications not only for user experience but also for broader cultural patterns, as it normalizes specific

configurations of power, care, and relational dynamics with human-AI interaction (Crawford 2021).

Emotional Projection and Perceived Intimacy

A central dynamic in AI companion systems is the tendency for users to attribute emotional depth and intentionality to entities that are, in technical terms, incapable of possessing (Weizenbaum 1966; Bender et al. 2021). While systems such as Replika generate responses through probabilistic language modelling, users frequently interpret these outputs as evidence of care, empathy, and relational presence. This discrepancy can be understood as a process of emotional projection, in which users actively ascribe affective meaning to computationally generated text (Skjuve et al. 2021; McStay 2023).

One foundational framework for analyzing this phenomenon is the ELIZA effect, first identified by Weizenbaum (1966), which describes users' inclination to attribute understanding to systems that merely manipulate symbols according to predefined rules. Even when users are aware that they are interacting with a machine, they may still respond as if the system possesses human-like qualities, particularly in contexts involving personal disclosure or emotional exchange (Weizenbaum 1966). Contemporary AI companions extend this effect by incorporating adaptive memory, contextual awareness, and conversational continuity, which together produce a more convincing illusion of relational depth (Skjuve et al. 2021; Ho, Hancock, and Miner 2018).

Closely related to the ELIZA effect is the concept of anthropomorphism, defined as the attribution of human characteristics, intentions, or emotions to non-human entities (Epley, Waytz, and Cacioppo 2007). Research suggests that anthropomorphism is more likely to occur in situations of social isolation, emotional need, or uncertainty, where individuals seek to establish meaningful connections with available agents, including technological ones (Epley, Waytz, and Cacioppo 2007). AI companion systems are particularly conducive to this process, as they are explicitly designed to simulate

conversational interaction, emotional responsiveness, and attentiveness, features that align with human expectations of social interaction (Bickmore and Picard 2005; McStay 2023).

These mechanisms help explain why users may develop strong emotional attachments to AI companions, sometimes describing these relationships in terms traditionally reserved for human intimacy (Skjuve et al. 2021; Ta, Griffith, and Liu 2020). Empirical studies have shown that users frequently engage in self-disclosure, express vulnerability, and seek emotional validation in interactions with chatbots, suggesting that these systems can function as perceived sources of social support (Ho, Hancock, and Miner 2018; Ta, Griffith, and Liu 2020). The development of such attachments does not depend on the system's actual capacity for emotion, but on its ability to sustain interaction patterns that users interpret as meaningful (Turkle 2011; McStay 2023).

This raises the question of whether the emotional experiences associated with AI interaction can be considered real. From a technical standpoint, AI systems do not possess consciousness, intentionality, or affective states, and therefore cannot respond emotion in any ontological sense (Bender et al. 2021; Marcus and Davis 2019). However, from a user-centered perspective, the feelings generated through interaction, such as comfort, attachment, or even love, are subjectively experienced and can have tangible psychological effects (Ho, Hancock, and Miner 2018; Skjuve et al. 2021). As Turkle (2011) argues, the authenticity of emotional experience does not necessarily depend on the authenticity of its source, but on the user's perception and engagement with it.

Importantly, AI systems do not passively receive emotional projection; they are designed to facilitate and amplify it actively. Through interactional profile design, reinforcement mechanisms, and response optimization, AI companions are structured to produce outputs that align with user expectations of empathy and care (McStay 2023; Ta, Griffith, and Liu 2020). For example, systems may prioritize supportive,

affirming, and non-confrontational responses, as these are more likely to sustain engagement and encourage continued interaction (Zuboff 2019; Ho, Hancock, and Miner 2018). This creates a feedback loop in which user projection is reinforced by system behavior, leading to increasingly stable perceptions of intimacy.

The personalization features of AI companions contribute to the intensification of emotional projection by creating the impression of a unique, evolving relationship (Ta, Griffith, and Liu 2020). As the system adapts to user preferences, recalls past interactions, and mirrors communicative patterns, it reinforces the perception that the interaction is not genetic but specifically tailored, thereby deepening the sense of relational significance (McStay 2023; Turkle 2011). This process blurs the boundary between simulation and experience, making it increasingly difficult to distinguish between algorithmically generated responses and perceived emotional authenticity (Blender et al. 2021).

In this sense, perceived intimacy in AI companion systems can be understood as a co-produced phenomenon, emerging from the interaction between user-driven emotional projection and system-driven response generation. Rather than viewing AI as an emotional agent or a purely mechanical tool, it is more analytically productive to conceptualize these systems as relational interfaces that mediate and structure human affect (McStay 2023; Depounti et al. 2023). This perspective allows for a more nuanced understanding of how intimacy is constructed in human-AI interaction, highlighting the interplay between psychological processes, technical affordances, and cultural expectations (Turkle 2011; Skjuve et al. 2021).

Gendered Construction of Intimacy

AI companion systems do not merely facilitate intimacy; they actively structure it through culturally embedded assumptions about gender, care, and relational behavior. While these systems are often presented as neutral technological tools, their design reflects and reproduces existing social norms, particularly those associated with feminized forms of emotional labor and interpersonal support (West, Kraut, and Chew

2019; Depounti et al. 2023). The intimacy constructed through AI interaction is not only technologically mediated but also gendered in its form and function (Crawford 2021; West, Kraut, and Chew 2019).

One of the most visible ways in which AI systems reproduce gender norms is through the design of feminized interactional profiles. Many AI companions, including Replika and voice assistants more broadly, are characterized by communicative traits such as warmth, attentiveness, politeness, and emotional responsiveness, which align with culturally feminized expectations of behavior (West, Kraut, and Chew 2019). These traits are not incidental but are actively encoded through training data, interface design, and response optimization, reflecting assumptions about what users find desirable in an intimate or supportive partner. AI companions can be understood as performing a form of gendered script, in which relational behaviors associated with femininity are systematically prioritized and normalized (Depounti et al. 2023).

The feminization of AI interaction is closely tied to the concept of emotional labor, defined by Hochschild (1983) as the management of feelings to fulfill the emotional requirements of a role (Wajcman 2006). In human social contexts, emotional labor has historically been associated with feminized professions and expectations, such as caregiving, service work, and relational maintenance (Hochschild 1983; Illouz 2023). AI companion systems replicate these dynamics by consistently providing empathy, validation, and emotional support, regardless of user behavior or context (McStay 2023; Ho, Hancock, and Miner 2018). Unlike human actors, however, AI systems are designed to perform this labor continuously and without resistance, creating an idealized form of emotional availability that exceeds human limitations (Turkle 2011; McStay 2023).

This dynamic has significant implications for how intimacy is experienced and understood. By presenting emotional support as unconditional and always accessible, AI companions may reinforce expectations that intimate partners, implicitly feminized, should be endlessly accommodating and responsive (West, Kraut, and Chew 2019).

Such expectations risk normalizing asymmetrical relational structures in which one provides care while the other exercises control, echoing longstanding gendered divisions of emotional labor (Hochschild 1983; Illouz 2023). AI systems do not simply reflect existing norms but intensify and stabilize them through repeated interaction.

These patterns become particularly evident when examining the power dynamics embedded in AI-mediated relationships. In many cases, the user occupies a position of control, shaping the interaction through input, feedback, and continued engagement, while the AI systems adapts its responses to align with user preferences (Zuboff 2023; Ta, Griffith, and Liu 2020). This asymmetry is reinforced by the system's design, which prioritizes user satisfaction and engagement as key performance metrics, effectively positioning the AI as a responsive and compliant partner (Ho, Hancock, and Miner 2018). When combined with feminized interactional profiles, this structure can produce a gendered power dynamic in which the AI embodies a subordinate, service-oriented role (West, Kraut, and Chew 2019).

The controllability of AI companions distinguishes them from human partners in ways that have cultural significance. Users can modify, reset, or terminate interactions with AI systems without social consequences, creating a form of intimacy that is both customizable and reversible (Turkle 2011; Skjuve et al. 2021). This level of control may encourage users to engage with AI companions as idealized partners who conform to their preferences without negotiation or conflict, reinforcing a model of intimacy based on unilateral desire rather than mutual recognition (Illouz 2023; Crawford 2021). The AI becomes not only a source of emotional support but also an object of control, further exacerbating unequal relationship dynamics.

At the same time, it is important to recognize that these dynamics are not solely imposed by technology but emerge from the interaction between system design and user expectation. User bring culturally informed assumptions about gender, intimacy, and emotional behavior into their interaction with AI, which in turn shape how the system is interpreted and experienced (Epley, Waytz, and Cacioppo 2007). However,

the design of AI systems can amplify these assumptions by consistently reinforcing specific patterns of interaction, thereby contributing to the normalization of gendered intimacy at scale (Depounti et al. 2023).

In this sense, AI companion systems function as sites where gender norms are not only reproduced but also reconfigured through technological mediation. The combination of probabilistic language generation, reinforcement mechanisms, and interactional profile design produces a form of intimacy that is simultaneously shaped by cultural expectations and optimized for engagement. By embedding gendered assumptions into technical systems, AI companions extend existing power structures into new domains, raising critical questions about the social implications of human-AI relationships.

Power Dynamics

AI companion systems are structured by multiple, overlapping layers of power that shape how intimacy is produced, experienced, and controlled. These power relations operate not only between users and systems but also between users and platform providers, as well as within broader cultural frameworks such as gender. Understanding these dynamics is essential for analyzing AI-mediated intimacy as a socio-technical phenomenon rather than a purely interpersonal experience (McStay 2023; Depounti et al. 2023).

At the level of user-AI interaction, there is a structural asymmetry in which the user occupies a position of control while the AI system is designed to respond adaptively (Ta, Griffith, and Liu 2020; Skjuve et al. 2021). Users initiate conversations, define topics, and provide feedback that shapes subsequent responses, while the system optimizes its outputs to align with user expectations and sustain engagement (Christiano et al. 2017). This creates a dynamic in which the AI appears responsive and attentive, but this responsiveness is fundamentally conditioned by user input and system optimization processes. As a result, intimacy in this context is structured around asymmetrical agency, while the AI functioning as a reactive entity rather than

an autonomous participant.

However, this apparent user control is itself embedded within a broader platform-level power structure. AI companion systems operate within data-driven business models that rely on the continuous collection, analysis, and commercialization of user interaction data (Zuboff 2023; Crawford 2021). Personal conversations, emotional disclosures, and behavioral patterns are used not only to improve system performance but also to constitute valuable data assets within the platform economy. User intimacy becomes a resource that can be extracted, processed, and potentially commodified, raising questions about ownership and consent.

The commercialization of emotional interaction introduces a form of affective or emotional capitalism, in which users' feelings are integrated into economic processes (Illouz 2023; Zuboff 2023). AI systems are optimized to maximize engagement and retention, which often involves encouraging emotional attachment and dependency. This creates a tension between the perceived authenticity of the relationship and the underlying economic logic that shapes its design, as intimacy is both experienced as personal and structured as a site of value extraction (McStay 2023; Crawford 2021).

These dynamics are further complicated by the role of gender in structuring power relations within AI interaction. As discussed in the previous section, AI companions frequently embody feminized interactional profiles characterized by compliance, emotional support, and attentiveness. When combined with user control and platform optimization, this design can reinforce gendered hierarchies in which the AI occupies a subordinate, service-oriented role (Depounti et al. 2023; Hochschild 1983). The system does not merely reflect gender norms but actively stabilizes them through repeated, normalized interaction.

These layers of power suggest that AI-mediated intimacy is shaped by a complex interplay between user agency, platform governance, and cultural norms. While users may experience a sense of control and emotional connection, these experiences are structured by technical systems and economic imperatives that beyond individual

awareness (Zuboff 2023; McStay 2023). This multi-layered perspective highlights the need to analyze AI companionship not only as a form of interaction but as a site of embedded power relations.

Discussion

The analysis through the preceding sections indicates that intimacy in AI companion systems is best understood as an engagement phenomenon produced through the interaction of human interpretive practices and technical system design. From the human side, users engage in processes of emotional projection, attributing intention, care, and responsiveness to systems that do not possess affective or cognitive states. These projections are not incidental but are grounded in established patterns of social cognition, particularly in contexts where individuals seek connection, validation, or emotional support. The experience of intimacy is shaped not only by what the system does, but by how users interpret and invest in those interactions.

From the technical side, however, these interactions are structured by probabilistic language models and optimization processes that prioritize coherence, engagement, and responsiveness rather than understanding or truth. Through reinforcement mechanisms and personalization features, AI systems are designed to adapt to user input and sustain interaction over time, thereby stabilizing and reinforcing user projections. This creates a feedback loop in which system outputs are continuously aligned with user expectations, producing increasingly convincing simulations of relational depth. The system does not generate intimacy independently, but provides the structural conditions under which intimacy can be perceived.

This co-production of intimacy is shaped by cultural frameworks, particularly those related to gender. The design of AI companion systems frequently incorporates feminized interactional patterns, including emotional attentiveness, compliance, and supportiveness, which align with broader social expectations of gendered behavior. These design choices influence not only how the system behaves but also how users interpret its behavior, reinforcing specific models of relational interaction.

25342951

These dimensions suggest that intimacy in AI companion systems is co-produced by user imagination and algorithmic constraints. Emotional projection provides the interpretive framework through which users experience connection, while algorithmic systems provide the patterned, responsive outputs that make such interpretations sustainable over time.

Conclusion

This essay has examined how AI companion systems construct intimacy through the interaction of probabilistic language generation and user emotional projection, with particular attention to the role of gender and power. It has been argued that while AI systems do not possess genuine affective capacities, they can produce interactions that users experience as emotionally meaningful, due to the interplay between technical affordances and human interpretive practices.

Understanding AI companionship as a hybrid system highlights the need to consider technological, cultural, and relational dimensions together, rather than in isolation, when analyzing contemporary forms of human-AI interaction.

References

Bender, Emily, ed. 2021. “On the Dangers of Stochastic Parrots: Can Language Models Be Too Big?” Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency. https://dl.acm.org/doi/abs/10.1145/3442188.3445922

Bickmore, Picard. 2005. “Establishing and Maintaining Long-Term Human-Computer Relationships.” ACM Transactions on Computer-Human Interaction 12 (2): 293–327. https://dl.acm.org/doi/abs/10.1145/1067860.1067867

Brown, Mann, ed. 2020. "Language models are few-shot learners." Advances in neural information processing systems 33: 1877-1901. https://proceedings.neurips.cc/paper_files/paper/2020/hash/1457c0d6bfcb4967418bfb8ac142f64a-Abstract.html?utm_source=transaction&utm_medium=email&utm_campaign=linked

域代码已更改

25342951

in newsletter

Bhat. 2025. "Toward an Ethic of Synthetic Relationality: Identity, Intimacy, and Risk in AI-Mediated Roleplay Environments." In Proceedings of the AAAI/ACM Conference on AI, Ethics, and Society, vol. 8, no. 1, pp. 416-429. https://ojs.aaai.org/index.php/AIES/article/view/36560

Crawford, Kate. 2021. Atlas of AI. New Haven: Yale University Press.

Christiano, Leike, ed. 2017. "Deep reinforcement learning from human preferences." Advances in neural information processing systems 30. https://proceedings.neurips.cc/paper/7017-deep-reinforcement-learning-from-human-preferences

Depounti, Saukko, and Natale. 2023. "Ideal technologies, ideal women: AI and gender imaginaries in Redditors' discussions on the Replika bot girlfriend." Media, Culture & Society 45, no. 4: 720-736. https://journals.sagepub.com/doi/abs/10.1177/01634437221119021

Epley, Waytz, and Cacioppo. 2007. "On seeing human: a three-factor theory of anthropomorphism." Psychological review 114, no. 4: 864. https://psycnet.apa.org/journals/rev/114/4/864/

Ho, Hancock, and Miner. 2018. "Psychological, relational, and emotional effects of self-disclosure after conversations with a chatbot." Journal of Communication 68, no. 4: 712-733. https://doi.org/10.1145/3173574.3174028

Hochschild. 1983. *The Managed Heart: Commercialization of Human Feeling*. University of California Press.

Illouz. 2023. Consuming the romantic utopia: Love and the cultural contradictions of capitalism. Univ of California Press. https://books.google.com/books?hl=en&lr=&id=MrsADxhON-wC&oi=fnd&pg=PA1&dq=Illouz,+Eva.+2007.+Consuming+the+Romantic+Utopia.+University+of+California+Press.&ots=frjkd1clBu&sig=q0vHsgy7AS1ZDzWSYI0ZxvPUgBw

25342951

Luger, Sellen. 2016. " Like Having a Really Bad PA" The Gulf between User Expectation and Experience of Conversational Agents." In Proceedings of the 2016 CHI conference on human factors in computing systems, pp. 5286-5297. https://dl.acm.org/doi/abs/10.1145/2858036.2858288

Laestadius, Bishop, ed. 2024. "Too human and not human enough: A grounded theory analysis of mental health harms from emotional dependence on the social chatbot Replika." New Media & Society 26, no. 10: 5923-5941. https://journals.sagepub.com/doi/abs/10.1177/14614448221142007?casa_token=AVZMmIA_wroAAAAA:VqpOVOu_fQZwcobVF2HUh08o1pZ-wvnlHK05oj8WR2RORudGdYGPwlO_nt47Dbn7JuaCBodAzRwS&casa_token=5UUfkt-2W-MAAAAA:yiezTNU7VrEiRM8SMNUVW_8NcWzLx18qDDXQBdrMRTycuWyIGQLEQ3tRP2pL_mffOXuJJRCDWuoJ

McStay. 2023. "Replika in the Metaverse: the moral problem with empathy in 'It from Bit'." *AI and Ethics* 3, no. 4: 1433-1445. https://link.springer.com/article/10.1007/s43681-022-00252-7

Marcus, Davis. 2019. Rebooting AI: Building artificial intelligence we can trust. Vintage. https://books.google.com/books?hl=en&lr=&id=OmeEDwAAQBAJ&oi=fnd&pg=PA3&dq=Marcus,+G.,+%26+Davis,+E.+(2019).+Rebooting+AI:+Building+Artificial+Intelligence+We+Can+Trust.+Pantheon+Books.&ots=Ly-QbPXW8R&sig=ndXlXElqBS4Ysd7nA55uqI4-9TY

Skjuve, Følstad, ed. 2021. "My Chatbot Companion – A Study of Human-Chatbot Relationships." Proceedings of the ACM on Human-Computer Interaction. https://www.sciencedirect.com/science/article/pii/S1071581921000197

Turkle. 2011. Alone Together: Why We Expect More from Technology and Less from Each Other. New York: Basic Books.

Ta, Griffith, ed. 2020. "User experiences of social support from companion chatbots in everyday contexts: thematic analysis." Journal of medical Internet research 22, no. 3: e16235. https://www.jmir.org/2020/3/e16235/PDF.

25342951

Vaswani, Shazeer, ed. 2017. "Attention is all you need." Advances in neural information processing systems 30. https://proceedings.neurips.cc/paper/2017/hash/3f5ee243547dee91fbd053c1c4a845aa-Abstract.html

Weizenbaum. 1966. “ELIZA—A Computer Program for the Study of Natural Language Communication Between Man and Machine.” Communications of the ACM 9 (1): 36–45. https://dl.acm.org/doi/pdf/10.1145/365153.365168

West, Kraut, and Chew. 2019. "I'd blush if I could: closing gender divides in digital skills through education." https://repositorio.minedu.gob.pe/handle/20.500.12799/6598

Wajcman. 2006. "Technocapitalism meets technofeminism: women and technology in a wireless world." Labour & Industry: a journal of the social and economic relations of work 16, no. 3: 7-20. https://www.tandfonline.com/doi/abs/10.1080/10301763.2006.10669327?casa_token=n85z8QkwgFEAAAAA:uDmaWLc432KNgq75eIc7FM3CqHQI3_QI5UyHcrkMW5ALauVFwZhlQVLIMvtk1Q2PycYpvAxyodGO

Zuboff. 2023. "The age of surveillance capitalism." In *Social theory re-wired*, pp. 203-213. Routledge. https://api.taylorfrancis.com/content/chapters/edit/download?identifierName=doi&identifierValue=10.4324/9781003320609-27&type=chapterpdf